

CS2040S

AY20/21 sem 2


github.com/jovynlt

ORDERS OF GROWTH

definitions


$$T(n) = O(f(n))$$

if $\exists c, n_0 > 0$ such that for all $n > n_0$, $T(n) \leq cf(n)$




$$T(n) = \Omega(f(n))$$

if $\exists c, n_0 > 0$ such that for all $n > n_0$, $T(n) \geq cf(n)$



$$T(n) = \Theta(f(n))$$

$$\Leftrightarrow T(n) = O(f(n)) \text{ and } T(n) = \Omega(f(n))$$



properties

Let $T(n) = O(f(n))$ and $S(n) = O(g(n))$

- addition: $T(n) + S(n) = O(f(n) + g(n))$
- multiplication: $T(n) * S(n) = O(f(n) * g(n))$
- composition: $f_1 \circ f_2 = O(g_1 \circ g_2)$
 - only if both functions are increasing
- if/else statements: $\text{cost} = \max(c_1, c_2) \leq c_1 + c_2$
- max: $\max(f(n), g(n)) \leq f(n) + g(n)$

notable

- $\sqrt{n} \log n$ is $O(n)$
- $O(2^{2n}) \neq O(2^n)$
- $O(\log(n!)) = O(n \log n) \rightarrow$ sterling's approximation

space complexity

- the maximum space incurred **at any time at any point**
- NOT the maximum space incurred altogether!
- assumption: once we exit the function, we release all memory that was used

SORTING

overview

- BubbleSort** - compare adjacent items and swap
- SelectionSort** - takes the smallest element, swaps into place
- InsertionSort** - from left to right: swap element leftwards until it's smaller than the next element. repeat for next element
 - tends to be faster than the other $O(n^2)$ algorithms
- MergeSort** - mergeSort first half; mergeSort second half; merge
 - partition algorithm: $O(n)$
 - first element as partition. 2 pointers from left to right
 - left pointer moves until element > pivot
 - right pointer moves until element < pivot
 - swap elements until left = right.
 - then swap partition and left=right index.
- QuickSort**

optimisations of QuickSort

- array of duplicates: $O(n^2)$ without 3-way partitioning
- stable if the partitioning algo is stable.
- extra memory allows quickSort to be stable.

choice of pivot

- worst case $O(n^2)$: first/last/middle element
- worst case $O(n \log n)$: median/random element
 - split by fractions: $O(n \log n)$
- choose at random: runtime is a random variable

quickSelect

- $O(n)$ - to find the k^{th} smallest element
- after partitioning, the partition is always in the correct position

TREES

binary search trees (BST)

- a BST is either empty, or a node pointing to 2 BSTs.
- tree balance depends on order of insertion
- balanced tree: $O(h) = O(\log n)$
- for a full-binary tree of size n , $\exists k \in \mathbb{Z}^+$ s.t. $n = 2^k - 1$

BST operations


- height**, $h(v) = \max(h(v.\text{left}), h(v.\text{right}))$
 - leaf nodes: $h(v) = 0$
- modifying operations**
 - search, insert** - $O(h)$
 - delete** - $O(h)$
 - case 1: no children - remove the node
 - case 2: 1 child - remove the node, connect parent to child
 - case 3: 2 children - delete the successor; replace node with successor
- query operations**
 - searchMin** - $O(h)$ - recurse into left subtree
 - searchMax** - $O(h)$ - recurse into right subtree
 - successor** - $O(h)$
 - if node has a right subtree: **searchMin**(v.right)
 - else: traverse upwards and return the first parent that contains the key in its left subtree

AVL Trees

height-balanced


- $\Leftrightarrow |v.\text{left}.\text{height} - v.\text{right}.\text{height}| \leq 1$
- each node is augmented with its height - $v.\text{height} = h(v)$
- space complexity: $O(LN)$ for N strings of length L

rebalancing




- insertion: max. 2 rotations
- deletion: recurse all the way up
- rotations can create every possible tree shape.


[case 1] B is **balanced**: right-rotate



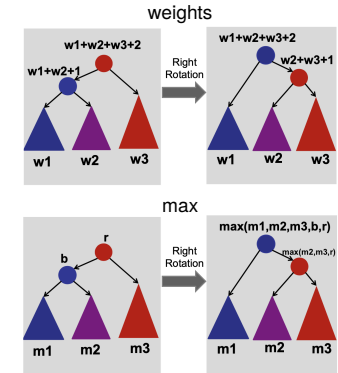
[case 2] B is **left-heavy**: right-rotate



[case 3] B is **right-heavy**: left-rotate(v.left), right-rotate(v)




updating nodes after rotation




Trie

- search, insert** - $O(L)$ (for string of length L)
- space: $O(\text{size of text} \times \text{overhead})$




interval trees



- search(key)** $\Rightarrow O(\log n)$
 - if value is in root interval, return
 - if value > max(left subtree), recurse right
 - else recurse left (go left only when can't go right)
- all-overlaps $\Rightarrow O(k \log n)$ for k overlapping intervals


kd-Tree



- stores geometric data (points in an (x, y) plane)
- alternates splitting (partitioning) via x and y coordinates
- construct**(points[]) $\Rightarrow O(n \log n)$
- search**(point) $\Rightarrow O(h)$
- searchMin**() $\Rightarrow T(n) = 2T(\frac{n}{4}) + O(1) \Rightarrow O(\sqrt{n})$


(a, b)-trees

e.g. a (2, 4)-tree storing 18 keys



- rules
 - (a, b) -child policy where $2 \leq a \leq (b + 1)/2$

	# keys		# children	
node type	min	max	min	max
root	1	$b - 1$	2	b
internal	$a - 1$	$b - 1$	a	b
leaf	$a - 1$	$b - 1$	0	0
 - an internal nodes has one more child than its number of keys
 - all leaf nodes must be at the same depth from the root
- terminology (for a node z)
 - key range - range of keys covered in subtree rooted at z
 - keylist - list of keys within z
 - treelist - list of z 's children
- max height = $O(\log_a n) + 1$
- min height = $O(\log_b n)$
- search(key) $\Rightarrow O(\log n)$
 - = $O(\log_2 b \cdot \log_a n)$ for binary search at each node
- insert(key) $\Rightarrow O(\log n)$
- split() a node with too many children
 - use median to split the keylist into 2 halves
 - move median key to parent; re-connect remaining nodes
 - (if the parent is now unbalanced, recurse upwards; if the root is reached, median key becomes the new root)



- delete(key) $\Rightarrow O(\log n)$
 - if the node becomes empty, merge(y, z) - join it with its left sibling & replace it with their parent
-
- if the combined nodes exceed max size: share(y, z) = merge(y, z) then split()

B-Tree

- $(B, 2B)$ -trees $\Rightarrow (a, b)$ -tree where $a = B, b = 2B$
- possible augmentation: use a LinkedList to connect between each level

Merkle Trees

- binary tree - nodes augmented with a hash of their children
- same root value = identical tree

HASH TABLES

- disadvantage: no successor/predecessor operation

hashing

- Let the m be the table size; let n be the number of items; let $cost(h)$ be the cost of the hash function
- load(hash table), $\alpha = \frac{n}{m}$
 - = average number of items per bucket
 - = expected number of items per bucket

hashing assumptions

- simple uniform hashing assumption
 - every key has an equal probability of being mapped to every bucket
 - keys are mapped independently
- uniform hashing assumption
 - every key is equally likely to be mapped to every permutation, independent of every other key.
 - NOT fulfilled by linear probing

properties of a good hash function

- able to enumerate all possible buckets - $h : U \rightarrow \{1..m\}$
 - for every bucket $j, \exists i$ such that $h(key, i) = j$
- simple uniform hashing assumption

hashCode

rules for the hashCode() method

- always returns the same value, if the object hasn't changed
- if two objects are equal, they return the same hashCode

rules for the equals method

- reflexive - `x.equals(x) => true`
- symmetric - `x.equals(y) => y.equals(x)`
- transitive - `x.equals(y), y.equals(z) => x.equals(z)`
- consistent - always returns the same answer
- null is null - `x.equals(null) => false`

chaining

- time complexity
 - insert(key, value) - $O(1 + cost(h)) \Rightarrow O(1)$
 - for n items: expected maximum cost
 - $= O(\log n)$
 - $= \Theta(\frac{\log n}{\log(\log(n))})$
 - search(key)
 - worst case: $O(n + cost(h)) \Rightarrow O(n)$
 - expected case: $O(\frac{n}{m} + cost(h)) \Rightarrow O(1)$
- total space: $O(m + n)$

open addressing - linear probing

- redefined hash function: $h(k, i) = h(k, 1) + i \bmod m$
- delete(key)
 - use a tombstone value - DON'T set to null
- performance
 - if the table is $\frac{1}{4}$ full, then there will be clusters of size $\Theta(\log n)$
 - expected cost of an operation, $E[\#probes] \leq \frac{1}{1-\alpha}$ (assume $\alpha < 1$ and uniform hashing)
- advantages
 - saves space (use empty slots vs linked list)
 - better cache performance (table is one place in memory)

- rarely allocate memory (no new list-node allocation)
- disadvantages
 - more sensitive to choice of hash function (clustering)
 - more sensitive to load (as $\alpha \rightarrow 1$, performance degrades)

double hashing

- for 2 functions f, g , define
- $$h(k, i) = f(k) + i \cdot g(k) \bmod m$$
- if $g(k)$ is relatively prime to m , then $h(k, i)$ hits all buckets
 - e.g. for $g(k) = n^k, n$ and m should be coprime.

table size

- assume chaining & simple uniform hashing
- let m_1 = size of the old hash table; m_2 = size of the new hash table; n = number of elements in the hash table
- growing the table: $O(m_1 + m_2 + n)$
 - rate of growth

table growth	resize	insert n items
increment by 1	$O(n)$	$O(n^2)$
double	$O(n)$	$O(n)$, average $O(1)$
square	$O(n^2)$	$O(n)$

PROBABILITY THEORY

- if an event occurs with probability p , the expected number of iterations needed for this event to occur is $\frac{1}{p}$.
- for random variables: expectation is always equal to the probability
- linearity of expectation: $E[A + B] = E[A] + E[B]$

UNIFORMLY RANDOM PERMUTATION

- for an array of n items, every of the $n!$ possible permutations are producible with probability of exactly $\frac{1}{n!}$
 - the number of outcomes should distribute over each permutation uniformly. (i.e. $\frac{\# \text{ of outcomes}}{\# \text{ of permutations}} \in \mathbb{N}$)
- probability of a specific item remaining in its initial position = $\frac{1}{n}$
- KnuthShuffle: for every element in array A , swap it with a random index in array A . $\Rightarrow O(n)$

sorting					
sort	best	average	worst	stable?	memory
bubble	$\Omega(n)$	$O(n^2)$	$O(n^2)$	✓	$O(1)$
selection	$\Omega(n^2)$	$O(n^2)$	$O(n^2)$	×	$O(1)$
insertion	$\Omega(n)$	$O(n^2)$	$O(n^2)$	✓	$O(1)$
merge	$\Omega(n \log n)$	$O(n \log n)$	$O(n \log n)$	✓	$O(n)$
quick	$\Omega(n \log n)$	$O(n \log n)$	$O(n^2)$	×	?

sorting invariants	
sort	invariant (after k iterations)
bubble	largest k elements are sorted
selection	smallest k elements are sorted
insertion	first k elements are in order
merge	—
quick	partition is in the right position

searching	
search	average
linear	$O(n)$
binary	$O(\log n)$
quickSelect	$O(n)$

data structures (search/insert) assuming $O(1)$ comparison cost

data structure	search	insert
sorted array	$O(\log n)$	$O(n)$
unsorted array	$O(n)$	$O(1)$
linked list	$O(n)$	$O(1)$
tree	$O(\log n)$ or $O(h)$	$O(\log n)$ or $O(h)$
trie	$O(L)$	$O(L)$
dictionary	$O(\log n)$	$O(\log n)$
symbol table	$O(1)$	$O(1)$
chaining	$O(n + cost(h))$	$O(1 + cost(h))$
open addressing	$O(1)$	$O(1)$

orders of growth

$$T(n) = 2T(\frac{n}{2}) + O(n) \Rightarrow O(n \log n)$$
$$T(n) = T(\frac{n}{2}) + O(n) \Rightarrow O(n)$$
$$T(n) = 2T(\frac{n}{2}) + O(1) \Rightarrow O(n)$$
$$T(n) = T(\frac{n}{2}) + O(1) \Rightarrow O(\log n)$$
$$T(n) = 2T(n - 1) + O(1) \Rightarrow O(2^n)$$
$$T(n) = 2T(\frac{n}{2}) + O(n \log n) \Rightarrow O(n(\log n)^2)$$
$$T(n) = 2T(\frac{n}{4}) + O(1) \Rightarrow O(\sqrt{n})$$
$$T(n) = T(n - c) + O(n) \Rightarrow O(n^2)$$